

Algorithmic Game Theory and Applications

Lecture 18:

Auctions and Mechanism Design II: a little social choice theory, the VCG Mechanism, and Market Equilibria

Kousha Etessami

Reminder: Food for Thought: sponsored search auctions

Question

How should advertisement slots on your Google search page be auctioned?

You may have already heard that Google uses a so-called “Generalized Second Price Auction” mechanism to do this.

But why do they do so? Is there any “better” way?

Before we can even begin to address such questions (which are in fact largely outside the scope of this course), it is useful to get a better understanding of the broader picture in economic theory within which [auction theory](#) and [mechanism design](#) operate.

Again, reference reading for today's lecture includes: Chapters 9, 11, 12, 13, & 28 of AGT book.

Also very useful are the relevant chapters of the excellent textbook: A. Mas-Colell, M. Whinston, & J. Green, *Microeconomic Theory*, 1995.

Walrasian Market Equilibrium

An Exchange Economy

- n agents, and m **divisible goods** (commodities, services, etc.).
- Each agent i starts with an initial **endowment** of goods,
 $w_i = (w_{i,1}, \dots, w_{i,m}) \in \mathbb{R}_{\geq 0}^m$.
- Each agent i has a **utility function**, $u_i(x)$, that defines how much it “likes” a given bundle $x \in \mathbb{R}_{\geq 0}^m$ of goods.

Assume utility functions satisfy certain “reasonable” conditions:

- (1) **Continuity**; (2) **Quasi-concavity**; (3) **Non-satiation**
- Assume agent endowments satisfy certain “reasonable conditions”.
E.g., each agent has a **positive endowment** of every good. (**Much less suffices**: e.g., each agent has something another agent “wants”, and the graph of this is strongly-connected.)
- Given a price vector, $p \in \mathbb{R}_{\geq 0}^m$, each agent i has an **optimal demand set**, $D_i(p) \subseteq \mathbb{R}_{\geq 0}^m$. This is the set of **bundles of goods**, x_i , that **maximize agent i 's utility**, and which i **can afford to buy** at prices p by selling its endowment w_i at prices p .

Background: existence of market equilibrium

Market (price) equilibrium for an exchange economy

A non-zero price vector, $p^* \geq 0$, together with bundles, x_i^* , of goods for each agent i , constitutes a **market (price) equilibrium** for an exchange economy, if at prices p^* , for every agent i , x_i^* is an **optimal demand bundle**, and moreover with these demands **the market clears** (i.e., basically **supply = demand**). In other words:

- For all agents i , $x_i^* \in D_i(p^*)$. (All agents optimize utility at prices p^*)
- $\sum_i x_i^* \leq \sum_i w_i$. (Demands do not exceed supply for any commodity.)
- Furthermore, for all goods j , if $(\sum_i x_i^*)_j < (\sum_i w_i)_j$, then $p_j^* = 0$.
(In other words, only **free goods** can have excess supply.)

Background: existence of market equilibrium

Market (price) equilibrium for an exchange economy

A non-zero price vector, $p^* \geq 0$, together with bundles, x_i^* , of goods for each agent i , constitutes a **market (price) equilibrium** for an exchange economy, if at prices p^* , for every agent i , x_i^* is an **optimal demand bundle**, and moreover with these demands **the market clears** (i.e., basically **supply = demand**). In other words:

- For all agents i , $x_i^* \in D_i(p^*)$. (All agents optimize utility at prices p^*)
- $\sum_i x_i^* \leq \sum_i w_i$. (Demands do not exceed supply for any commodity.)
- Furthermore, for all goods j , if $(\sum_i x_i^*)_j < (\sum_i w_i)_j$, then $p_j^* = 0$.
(In other words, only **free goods** can have excess supply.)

Theorem ([Arrow-Debreu, 1954])

Every exchange economy has a market (price) equilibrium.

Background: existence of market equilibrium

Market (price) equilibrium for an exchange economy

A non-zero price vector, $p^* \geq 0$, together with bundles, x_i^* , of goods for each agent i , constitutes a **market (price) equilibrium** for an exchange economy, if at prices p^* , for every agent i , x_i^* is an **optimal demand bundle**, and moreover with these demands **the market clears** (i.e., basically **supply = demand**). In other words:

- For all agents i , $x_i^* \in D_i(p^*)$. (All agents optimize utility at prices p^*)
- $\sum_i x_i^* \leq \sum_i w_i$. (Demands do not exceed supply for any commodity.)
- Furthermore, for all goods j , if $(\sum_i x_i^*)_j < (\sum_i w_i)_j$, then $p_j^* = 0$.
(In other words, only **free goods** can have excess supply.)

Theorem ([Arrow-Debreu, 1954])

Every exchange economy has a market (price) equilibrium.

Their proof is non-algorithmic. It crucially uses a **fixed point theorem**.

two quotes

Kamal Jain (MSR), 2006: “If your laptop can't find it, then neither can the market.”

Mas-Colell, Whinston, & Green, 1995 (standard graduate text in Microeconomic Theory): “A characteristic feature [of] economics is that for us the equations of equilibrium constitute the center of our discipline. By contrast, other sciences put more emphasis on the dynamic laws of change. The reason... is that economists are good at recognizing a state of equilibrium, but are poor at predicting precisely how an economy in disequilibrium will evolve...”

Or, to paraphrase: “Modern economic theory is largely non-algorithmic.” By contrast, **computer science** is very good at “dynamics” (algorithmics). A sizable portion of **Algorithmic Game Theory** research, broadly speaking, **aims to remedy this deficiency of modern economic theory.**

A side question: Do we really need money?

A side question: Do we really need money?

- Note that **money** is itself extraneous to the goals of agents in an exchange economy. Why do we need it?
- Couldn't we achieve the same “goals” of the market without money?
- In other words, couldn't we organize a “**barter economy**” to conduct a simultaneous exchange of goods between all agents, in a way that achieves **Pareto optimal** bundles of goods for every agent?

Recall: Pareto optimal means that there is no other reallocation of everyone's goods in which everybody's utility is at least as high, and somebody's utility is strictly higher.

Note: the **The First Fundamental Theorem of Welfare Economics** says that allocations of goods in any market equilibrium are Pareto optimal.

A side question: Do we really need money?

- Note that **money** is itself extraneous to the goals of agents in an exchange economy. Why do we need it?
- Couldn't we achieve the same “goals” of the market without money?
- In other words, couldn't we organize a “**barter economy**” to conduct a simultaneous exchange of goods between all agents, in a way that achieves **Pareto optimal** bundles of goods for every agent?

Recall: Pareto optimal means that there is no other reallocation of everyone's goods in which everybody's utility is at least as high, and somebody's utility is strictly higher.

Note: the **The First Fundamental Theorem of Welfare Economics** says that allocations of goods in any market equilibrium are Pareto optimal.

It is intuitively clear that such an exchange is **difficult** to do without money. We shall see that some things become **impossible** to do without money.

Background: Impossibility theorems in social choice theory

- Let C be a set of **outcomes**, and let L be the **set of total orderings** on C . A **social choice function** is a function $f : L^n \rightarrow C$.
- f can be **strategically manipulated** by player (**voter**) i , if for some $\prec_1, \dots, \prec_n \in L$, we have $c \prec_i c'$, and $c = f(\prec_1, \dots, \prec_n)$, and there exists some other ordering \prec'_i such that $c' = f(\prec_1, \dots, \prec'_i, \dots, \prec_n)$.
- f is called **incentive compatible (strategy-proof)** if it can not be strategically manipulated by anyone.
- f is a **dictatorship** if there is some player i such that for all \prec_1, \dots, \prec_n , $f(\prec_1, \dots, \prec_n) = c_i$, where c_i is the maximum outcome for player i , i.e., $\forall b \neq c_i, b \prec_i c_i$.

Theorem [Gibbard-Satterthwaite'73,'75]

If $|C| \geq 3$, and $f : L^n \mapsto C$ is an incentive compatible social choice function which is onto C (i.e., all outcomes are possible), then f is a dictatorship.

The Gibbard-Satterthwaite Theorem is closely related to the following famous theorem:

Arrow's Impossibility Theorem

There is no “really good” way to aggregate people's preference orders (over ≥ 3 alternatives) into a societal preference order.

“Really good” here means it should satisfy the following three criteria:

- *Unanimity*: if everybody has the same preference order, then that should become the societal preference order.
- *Independence of irrelevant alternatives*: The societal preference between any pair of outcomes a and b should depend only on the preference between a and b of all individuals, and not on their preferences for other alternatives $c \neq a, b$.
- *Non-dictatorship*: no individual's preferences should dictate the societies preferences.

Let's bring money back into the picture

- Let V be the set of “voters”.
- Suppose that instead of a “preference ordering”, each voter $i \in V$ has a (**non-negative**) “**monetary value function**”: $v_i : C \rightarrow \mathbb{R}_{\geq 0}$, where $v_i(c)$ says **how much money** a win by candidate c is **worth** for voter i .
- Suppose we want to **choose a candidate c^* that maximizes the total value for the entire voter population**, i.e., we want to choose:

$$f(v_1, \dots, v_n) = c^* \in \arg \max_c \sum_{i \in V} v_i(c)$$

- But voters could be **dishonest** and **lie** about their true value function. (Voter i can declare a different non-negative valuation function v'_i .)
- *Question:* **Can we incentivize the voters to tell the truth?**

Let's bring money back into the picture

- Let V be the set of “voters”.
- Suppose that instead of a “preference ordering”, each voter $i \in V$ has a (**non-negative**) “**monetary value function**”: $v_i : C \rightarrow \mathbb{R}_{\geq 0}$, where $v_i(c)$ says **how much money** a win by candidate c is **worth** for voter i .
- Suppose we want to **choose a candidate c^* that maximizes the total value for the entire voter population**, i.e., we want to choose:

$$f(v_1, \dots, v_n) = c^* \in \arg \max_c \sum_{i \in V} v_i(c)$$

- But voters could be **dishonest** and **lie** about their true value function. (Voter i can declare a different non-negative valuation function v'_i .)
- *Question:* **Can we incentivize the voters to tell the truth?**
- **Yes**, if we are allowed to ask them to **pay money**!

The Vickrey-Clarke-Groves (VCG) Mechanism

- Each voter $i \in V$ is asked to submit their nonnegative valuation function, v'_i . (v'_i may or may not be i 's **actual** valuation function v_i .)
- The mechanism then computes an optimal candidate, $f(v'_1, \dots, v'_n) := c^*$, that maximizes total value:

$$f(v'_1, \dots, v'_n) := c^* \in \arg \max_c \sum_{k \in V} v'_k(c)$$

This candidate c^* is **chosen as the winner** of the election.

- Moreover, each voter i has to **pay** an amount $p_i(c^*)$ to the mechanism, which is **independent of v'_i** , defined by:

$$p_i(c^*) := \left(\max_{c' \in C} \sum_{j \in V \setminus \{i\}} v'_j(c') \right) - \sum_{j \in V \setminus \{i\}} v'_j(c^*)$$

Key Point: These payments align the incentives of all voters: each voter i , even knowing the valuation functions v'_j declared by voters $j \neq i$, will want to declare a function v'_i that yields a winner $c^* = f(v'_1, \dots, v'_n)$ which maximizes $v_i(c^*) + \sum_{k \in V \setminus \{i\}} v'_k(c^*)$. So, it should declare $v'_i := v_i$.

Proposition

Let $f(v'_1, \dots, v'_n) = c^*$ be the outcome chosen as the winner by the VCG mechanism, based on their declared (**non-negative**) valuation functions v'_1, \dots, v'_n . Then **for every voter i** both its **payment $p_i(c^*)$** and its **purported utility $u'_i(c^*) = v'_i(c^*) - p_i(c^*)$** are **non-negative**, i.e.:

$$1.) \quad p_i(c^*) \geq 0 \quad \text{and} \quad 2.) \quad u'_i(c^*) := v'_i(c^*) - p_i(c^*) \geq 0$$

Proof: 1.) $p_i(c^*) = \left(\max_{c' \in C} \sum_{j \in V \setminus \{i\}} v'_j(c') \right) - \sum_{j \in V \setminus \{i\}} v'_j(c^*) \geq 0.$

$$\begin{aligned} 2.) \quad v'_i(c^*) - p_i(c^*) &= \left(\sum_{j \in V} v'_j(c^*) \right) - \left(\max_{c' \in C} \sum_{j \in V \setminus \{i\}} v'_j(c') \right) \\ &= \left(\max_{c \in C} \sum_{i \in V} v'_i(c) \right) - \left(\max_{c' \in C} \sum_{j \in V \setminus \{i\}} v'_j(c') \right) \geq 0 \end{aligned}$$

(The last ≥ 0 holds because we require $v'_j(c) \geq 0$ for all j and all $c \in C$.)

VCG is incentive-compatible (i.e., strategy-proof)

Theorem ([Vickrey, 1961],[Clarke, 1971], [Groves,1973])

The VCG mechanism is *incentive compatible* (i.e., strategy proof).
In other words, declaring their true valuation function v_i is a (weakly) dominant strategy for all players i .

Proof: Suppose players declare valuations v'_1, \dots, v'_n . Suppose player i 's true valuation is v_i . Let $c^* = f(v_i, v'_-i)$, and let $c'' = f(v'_i, v'_-i)$. Recall $c^* \in \arg \max_c v_i(c) + \sum_{k \in V \setminus \{i\}} v'_k(c)$. In other words, for all $c \in C$,
 $v_i(c^*) + \sum_{k \in V \setminus \{i\}} v'_k(c^*) \geq v_i(c) + \sum_{k \in V \setminus \{i\}} v'_k(c)$. Thus, $u_i(c^*) =$

$$\begin{aligned} v_i(c^*) - p_i(c^*) &= v_i(c^*) + \left(\sum_{k \in V \setminus \{i\}} v'_k(c^*) \right) - \left(\max_{c' \in C} \sum_{j \in V \setminus \{i\}} v'_j(c') \right) \\ &\geq v_i(c'') + \left(\sum_{k \in V \setminus \{i\}} v'_k(c'') \right) - \left(\max_{c' \in C} \sum_{j \in V \setminus \{i\}} v'_j(c') \right) \\ &= v_i(c'') - p_i(c'') = u_i(c''). \quad \square \end{aligned}$$

Example: Vickrey's second-price sealed-bid auction

- **Recall:** There is a single item being auctioned.
- Each bidder, i , declares its own monetary value (price), b_i , for getting the item.
- The highest bidder, say i^* , gets the item, and pays the **second highest bid price**, say $p^* = b_j$.
- The overall utility for bidder i^* who got the item is $v_{i^*} - p^*$. Both the valuation and payment (and thus utility) for all others who don't get it is 0.
- **Recall: Theorem:** the Vickrey auction is incentive compatible: every bidder i 's dominant strategy is to declare its true valuation v_i .
- In fact, **this is precisely the VCG mechanism**, applied in the setting of a single-item auction, where the set of **"candidates"**, now **"outcomes"**, is given by **"who gets the item?"**.
The mechanism makes sure that the item goes to the person who values it most, because that maximizes $\sum_i v_i(\text{outcome})$.

An ascending price auction closely related to Vickrey's

- In practice, Vickrey's sealed-bid auction are often not used.
- There are a variety of (good) reasons for this. (For example, bidder valuation may not be independent. This leads, e.g., to **Winner's curse**. Thus **information** about other people's valuation matters a lot.)
- Instead, often an **Ascending price (English) auction** is used:
 - The price starts at a low **reserve** price, and is repeatedly incremented by a small amount, $\epsilon > 0$, until all but one of the bidders drops out saying "I don't want it at that high a price".
 - At that point, the one remaining bidder gets the item, at the last price.
- It is intuitively clear that if we model these auctions as (Bayesian) games with **private valuations** then the outcome of the English auction should be "essentially the same" as that of the Vickrey auction.
- Things change a lot if we model them with **interdependent valuations**. (A model we haven't discussed: players don't know exactly their own valuation. Their valuation depends on everybody's private signal.). That is one reason why an English auction is often preferred.

Why not use VCG for multi-item auctions?

- Suppose there is a set J of items, with $|J| = k > 1$, up for auction.
- Suppose each bidder i is “single-minded” and values only one subset $J_i \subseteq J$ of the items, so he has a valuation $v_i(J') = v_i^* > 0$ for any $J' \supseteq J_i$, and $v_i(J') = 0$ for all $J' \not\supseteq J_i$. (So, we can specify valuation v_i by just giving (v_i^*, J_i) .)
- **Note:** if $J_i \cap J_{i'} \neq \emptyset$, then auction can't satisfy both bidders i and i' .
- Suppose, we want the auction outcome to maximize total value for all bidders.
- Why not use the VCG mechanism?
- A problem: for VCG we have to compute an outcome, c^* , that maximizes the total value $\sum_i v_i(c^*)$.
- **This is NP-hard in general, even when we always have $v_i^* = 1$:** easily encodes **maximum independent set**.

Proof of NP-hardness

Here is the simple reduction from the **maximum independent set** problem:

Let $G = (V, E)$ be any graph, with $V = \{s_1, \dots, s_n\}$.

We associate with each edge $e \in E$ an item in the auction.

We associate a player (bidder), i , with every vertex $s_i \in V$.

We let J_i be the set of edges adjacent to vertex v_i .

We let $v_i^* = 1$ for all i (i.e., every player has value 1 for its desired set).

Then finding an outcome c^* of the auction (i.e., a collection of disjoint sets J_i , indicating those bidders who get their desired set) such that this outcome maximizes $\sum_i v_i(c^*)$ is equivalent to finding a maximum independent set in the graph G .

Algorithmic Mechanism Design

- So, achieving incentive compatibility by means of VCG may involve solving NP-hard problems.
- Why not try the standard **approximation** approach used for coping with NP-hardness?

Problem: if the mechanism outputs only an approximate solution, and yet tries to compute VCG-like prices for it, then incentive compatibility totally breaks down.

- In some auction settings (like **single-minded bidders**), it is possible to design non-VCG mechanisms which are: (1) incentive compatible, (2) where prices and allocations are polynomial-time computable, and (3) such that the outcome yields a total valuation within the best possible approximation ratio of optimal (any better approximation ratio would be NP-hard).
- These are the kinds of problems addressed by *Algorithmic Mechanism Design*.

- There are other multi-item auction settings where the (exact) VCG mechanism remains useful and can be efficiently computed.
- One important and beautiful example of this is [Matching markets \(multi-item unit-demand auctions\)](#).

The setting of matching markets also gets us much closer to the setting of [sponsored search auctions](#) (also known as [position auctions](#)). These can be viewed as special subcases of matching markets and unit-demand auctions.

- We will discuss these next time.
- We will then end our lectures by presenting a general formal framework for Mechanism Design.